



TabPFN diagnostic: foundation-model baseline vs tuned ML + Putirka, opx pipelines

Diagnostic scope only: this figure sits outside the tuned-baseline bias-correction narrative (Core_05-07) and shows TabPFN for context.
Benchmark set: ExPetDB opx held-out test partition (citation-grouped, pre-registered regime edges; bootstrap 95% CI)

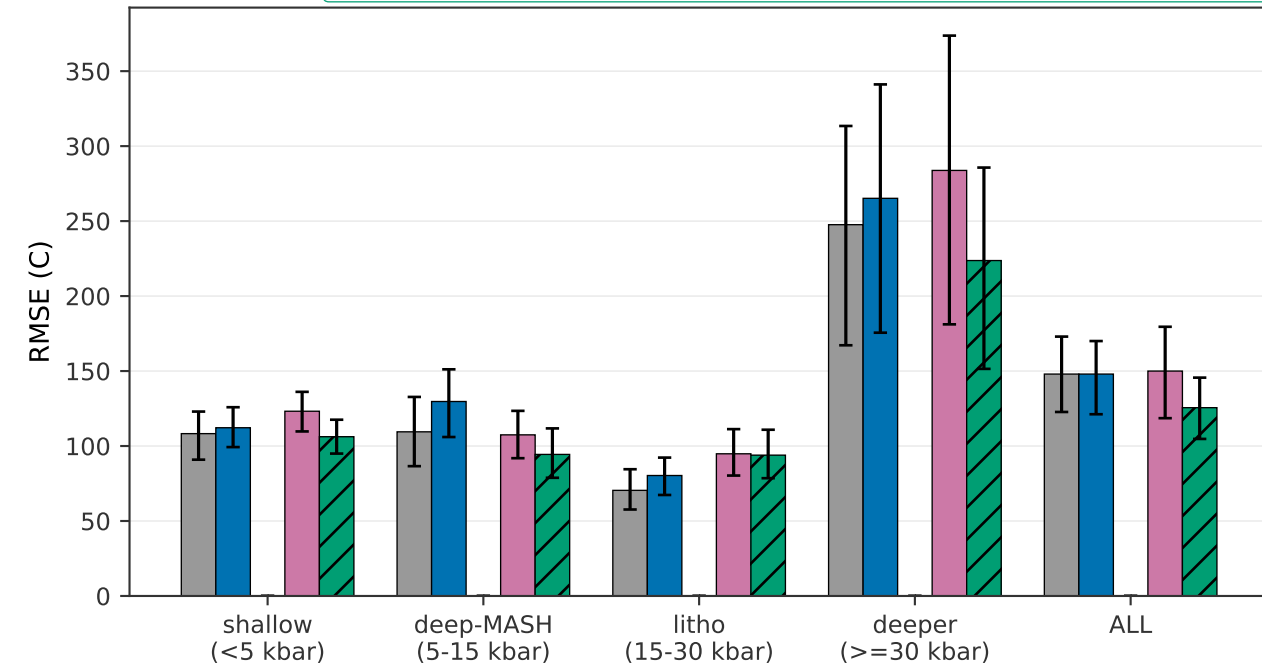
(a) Opx + Liquid T (C) ALL lowest RMSE (diagnostic): external [Putirka 28a]



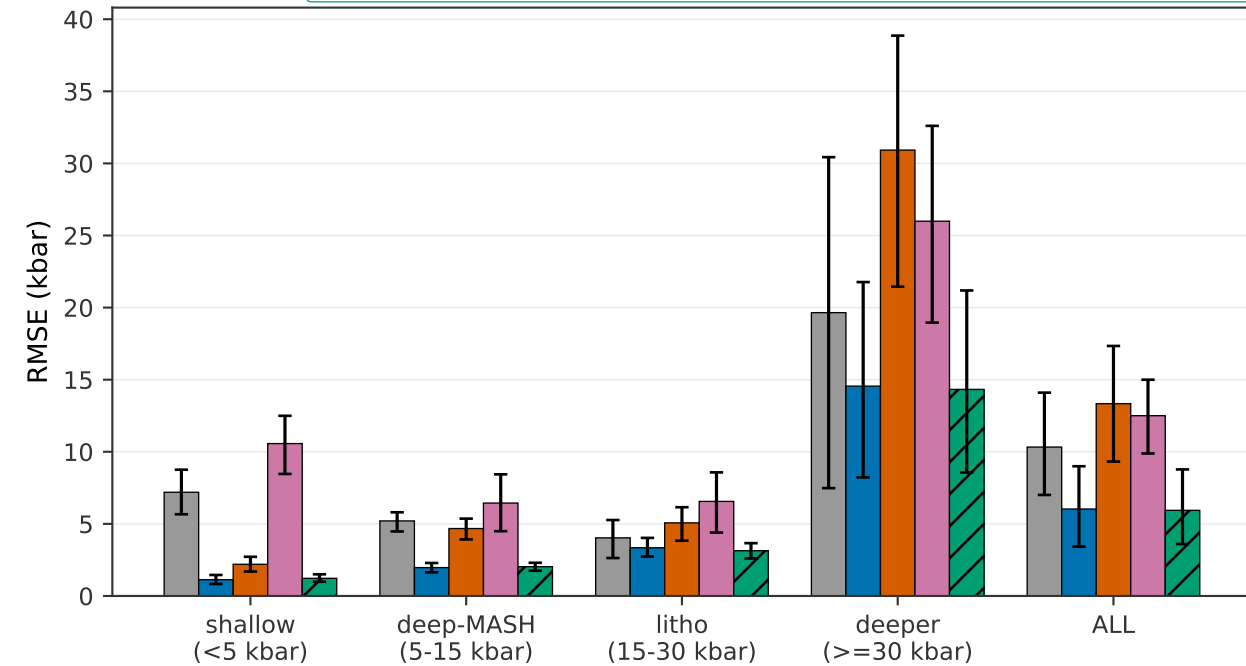
(b) Opx + Liquid P (kbar) ALL lowest RMSE (diagnostic): v10_corrected [MLP/raw]



(c) Opx only T (C) ALL lowest RMSE (diagnostic): tabpfn_corrected [TabPFN/raw (post)]



(d) Opx only P (kbar) ALL lowest RMSE (diagnostic): tabpfn_corrected [TabPFN/raw (post)]



v10 pre (tuned, uncorrected)
 v10 post (shipped form)
 Putirka (2008)
 TabPFN pre
 TabPFN post (shipped form)